

data-wrangling

June 9, 2020

```
<a href="https://coc1.us/corsera_da0101en_notebook_top">  
    
</a>
```

Data Analysis with Python

Data Wrangling

Welcome!

By the end of this notebook, you will have learned the basics of Data Wrangling!

Table of content

Identify and handle missing values

Identify missing values

Deal with missing values

Correct data format

```
</li>  
<li><a href="#data_standardization">Data standardization</a></li>  
<li><a href="#data_normalization">Data Normalization (centering/scaling)</a></li>  
<li><a href="#binning">Binning</a></li>  
<li><a href="#indicator">Indicator variable</a></li>
```

Estimated Time Needed: 30 min

What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

What is the fuel consumption (L/100k) rate for the diesel car?

Import data

You can find the “Automobile Data Set” from the following link:
<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will be using this data set throughout this course.

Import pandas

```
[45]: import pandas as pd  
import matplotlib.pyplot as plt
```

Reading the data set from the URL and adding the related headers.

URL of the dataset

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[46]: filename = "https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
      ↪CognitiveClass/DA0101EN/auto.csv"
```

Python list headers containing name of headers

```
[47]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
      ↪"num-of-doors", "body-style",
      ↪"drive-wheels", "engine-location", "wheel-base",
      ↪"length", "width", "height", "curb-weight", "engine-type",
      ↪"num-of-cylinders",
      ↪"engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
      ↪"peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter “names” equal to the Python list “headers”.

```
[48]: df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
[49]: # To see what the data set looks like, we'll use the head() method.
      df.head()
```

```
[49]:  symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0         3             ?  alfa-romero    gas      std         two
1         3             ?  alfa-romero    gas      std         two
2         1             ?  alfa-romero    gas      std         two
3         2          164      audi      gas      std         four
4         2          164      audi      gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi  3.47   2.68           9.0          111      5000      21
1      mpfi  3.47   2.68           9.0          111      5000      21
2      mpfi  2.68   3.47           9.0          154      5000      19
3      mpfi  3.19   3.40          10.0          102      5500      24
4      mpfi  3.19   3.40           8.0          115      5500      18
```

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data:

identify missing data

deal with missing data

correct data format

Identify and handle missing values

Identify missing values

Convert “?” to NaN

In the car dataset, missing data comes with the question mark “?”. We replace “?” with NaN (Not a Number), which is Python’s default missing value marker, for reasons of computational speed and convenience. Here we use the function:

to replace A by B

```
[50]: import numpy as np

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)
```

```
[50]:   symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero      gas      std         two
1         3             NaN  alfa-romero      gas      std         two
2         1             NaN  alfa-romero      gas      std         two
3         2            164      audi      gas      std         four
4         2            164      audi      gas      std         four

   body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...        130
1  convertible         rwd         front      88.6  ...        130
```

2	hatchback	rwd	front	94.5	...	152
3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

identify__missing__values

Evaluating for Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. There are two methods to detect missing data:

.isnull()

.notnull()

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
[51]: missing_data = df.isnull()
missing_data.head(5)
```

```
[51]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	False	True	False	False	False	False	
1	False	True	False	False	False	False	
2	False	True	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	False	False	False	False	...	False	
1	False	False	False	False	...	False	
2	False	False	False	False	...	False	
3	False	False	False	False	...	False	
4	False	False	False	False	...	False	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

“True” stands for missing value, while “False” stands for not missing value.

Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, “True” represents a missing value, “False” means the value is present in the dataset. In the body of the for loop the method “.value_counts()” counts the number of “True” values.

```
[52]: for column in missing_data.columns.values.tolist():
      print(column)
      print(missing_data[column].value_counts())
      print("")
```

symboling

False 205

Name: symboling, dtype: int64

normalized-losses

False 164

True 41

Name: normalized-losses, dtype: int64

make

False 205

Name: make, dtype: int64

fuel-type

False 205

Name: fuel-type, dtype: int64

aspiration

```

False      205
Name: aspiration, dtype: int64

num-of-doors
False      203
True        2
Name: num-of-doors, dtype: int64

body-style
False      205
Name: body-style, dtype: int64

drive-wheels
False      205
Name: drive-wheels, dtype: int64

engine-location
False      205
Name: engine-location, dtype: int64

wheel-base
False      205
Name: wheel-base, dtype: int64

length
False      205
Name: length, dtype: int64

width
False      205
Name: width, dtype: int64

height
False      205
Name: height, dtype: int64

curb-weight
False      205
Name: curb-weight, dtype: int64

engine-type
False      205
Name: engine-type, dtype: int64

num-of-cylinders
False      205
Name: num-of-cylinders, dtype: int64

```

```

engine-size
False      205
Name: engine-size, dtype: int64

fuel-system
False      205
Name: fuel-system, dtype: int64

bore
False      201
True        4
Name: bore, dtype: int64

stroke
False      201
True        4
Name: stroke, dtype: int64

compression-ratio
False      205
Name: compression-ratio, dtype: int64

horsepower
False      203
True        2
Name: horsepower, dtype: int64

peak-rpm
False      203
True        2
Name: peak-rpm, dtype: int64

city-mpg
False      205
Name: city-mpg, dtype: int64

highway-mpg
False      205
Name: highway-mpg, dtype: int64

price
False      201
True        4
Name: price, dtype: int64

```

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

“normalized-losses”: 41 missing data

“num-of-doors”: 2 missing data

“bore”: 4 missing data

“stroke” : 4 missing data

“horsepower”: 2 missing data

“peak-rpm”: 2 missing data

“price”: 4 missing data

Deal with missing data

How to deal with missing data?

drop data a. drop the whole row b. drop the whole column

replace data a. replace it by mean b. replace it by frequency c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

“normalized-losses”: 41 missing data, replace them with mean

“stroke”: 4 missing data, replace them with mean

“bore”: 4 missing data, replace them with mean

“horsepower”: 2 missing data, replace them with mean

“peak-rpm”: 2 missing data, replace them with mean

Replace by frequency:

“num-of-doors”: 2 missing data, replace them with “four”.

Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

Drop the whole row:

“price”: 4 missing data, simply delete the whole row

Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

Calculate the average of the column

```
[53]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```


Average of normalized-losses: 122.0

Replace “NaN” by mean value in “normalized-losses” column

```
[54]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

Calculate the mean value for ‘bore’ column

```
[55]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

Replace NaN by mean value

```
[56]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

Question #1:

According to the example above, replace NaN in “stroke” column by mean.

```
[57]: # Write your code below and press Shift+Enter to execute
      avg_stroke = df['stroke'].astype('float').mean(axis=0)
      df['stroke'].replace(np.nan, avg_stroke, inplace=True)
```

Double-click here for the solution.

Calculate the mean value for the ‘horsepower’ column:

```
[58]: avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
      print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

Replace “NaN” by mean value:

```
[59]: df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

Calculate the mean value for ‘peak-rpm’ column:

```
[60]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
      print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

Replace NaN by mean value:

```
[61]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the “value_counts()” method:

```
[62]: df['num-of-doors'].value_counts()
```

```
[62]: four      114
      two       89
      Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the “`idxmax()`” method to calculate for us the most common type automatically:

```
[63]: df['num-of-doors'].value_counts().idxmax()
```

```
[63]: 'four'
```

The replacement procedure is very similar to what we have seen previously

```
[64]: #replace the missing 'num-of-doors' values by the most frequent
      df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let’s drop all rows that do not have price data:

```
[65]: # simply drop whole row with NaN in "price" column
      df.dropna(subset=["price"], axis=0, inplace=True)

      # reset index, because we dropped two rows
      df.reset_index(drop=True, inplace=True)
```

```
[66]: df.head()
```

```
[66]:   symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0          3           122  alfa-romero    gas      std         two
1          3           122  alfa-romero    gas      std         two
2          1           122  alfa-romero    gas      std         two
3          2           164      audi    gas      std         four
4          2           164      audi    gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...      130
1  convertible         rwd         front      88.6  ...      130
2   hatchback         rwd         front      94.5  ...      152
3      sedan         fwd         front      99.8  ...      109
4      sedan         4wd         front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi  3.47    2.68           9.0        111      5000      21
1      mpfi  3.47    2.68           9.0        111      5000      21
2      mpfi  2.68    3.47           9.0        154      5000      19
3      mpfi  3.19    3.40          10.0        102      5500      24
4      mpfi  3.19    3.40           8.0        115      5500      18

      highway-mpg  price
```

```
0          27  13495
1          27  16500
2          26  16500
3          30  13950
4          22  17450
```

```
[5 rows x 26 columns]
```

Good! Now, we obtain the dataset with no missing values.

Correct data format

We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

`.dtype()` to check the data type

`.astype()` to change the data type

Lets list the data types for each column

```
[67]: df.dtypes
```

```
[67]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                 float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   object
stroke                 object
compression-ratio      float64
horsepower             object
peak-rpm               object
city-mpg               int64
```

```
highway-mpg      int64
price            object
dtype: object
```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

Convert data types to proper format

```
[68]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

Let us list the columns after the conversion

```
[69]: df.dtypes
```

```
[69]: symboling      int64
normalized-losses  int64
make              object
fuel-type         object
aspiration        object
num-of-doors      object
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length           float64
width            float64
height           float64
curb-weight       int64
engine-type       object
num-of-cylinders  object
engine-size       int64
fuel-system       object
bore              float64
stroke            float64
compression-ratio float64
horsepower        object
peak-rpm          float64
city-mpg          int64
highway-mpg       int64
price            float64
dtype: object
```

Wonderful!

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns “city-mpg” and “highway-mpg” are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard

We will need to apply data transformation to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

```
[70]: df.head()
```

```
[70]:   symboling  normalized-losses      make fuel-type aspiration \
0          3             122  alfa-romero      gas      std
1          3             122  alfa-romero      gas      std
2          1             122  alfa-romero      gas      std
3          2             164      audi      gas      std
4          2             164      audi      gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
3          four      sedan      fwd      front      99.8  ...
4          four      sedan      4wd      front      99.4  ...

   engine-size  fuel-system  bore  stroke  compression-ratio  horsepower  \
0          130      mpfi  3.47   2.68             9.0          111
1          130      mpfi  3.47   2.68             9.0          111
2          152      mpfi  2.68   3.47             9.0          154
3          109      mpfi  3.19   3.40            10.0          102
4          136      mpfi  3.19   3.40             8.0          115
```

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

[5 rows x 26 columns]

```
[71]: # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

```
[71]:      symboling  normalized-losses      make fuel-type aspiration \
0          3           122  alfa-romero      gas      std
1          3           122  alfa-romero      gas      std
2          1           122  alfa-romero      gas      std
3          2           164      audi      gas      std
4          2           164      audi      gas      std
```

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000.0	21	
1	mpfi	3.47	2.68	9.0	111	5000.0	21	
2	mpfi	2.68	3.47	9.0	154	5000.0	19	
3	mpfi	3.19	3.40	10.0	102	5500.0	24	
4	mpfi	3.19	3.40	8.0	115	5500.0	18	

	highway-mpg	price	city-L/100km
0	27	13495.0	11.190476
1	27	16500.0	11.190476
2	26	16500.0	12.368421
3	30	13950.0	9.791667
4	22	17450.0	13.055556

[5 rows x 27 columns]

Question #2:

According to the example above, transform mpg to L/100km in the column of “highway-mpg”, and change the name of column to “highway-L/100km”.

```
[72]: # Write your code below and press Shift+Enter to execute
df.rename(columns={'highway-mpg': 'high-L'}, inplace=True)
df.head()
```

```
[72]:   symboling  normalized-losses      make fuel-type aspiration \
0         3           122  alfa-romero    gas      std
1         3           122  alfa-romero    gas      std
2         1           122  alfa-romero    gas      std
3         2           164      audi    gas      std
4         2           164      audi    gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four     sedan      fwd      front      99.8  ...
4         four     sedan      4wd      front      99.4  ...

   fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0         mpfi  3.47   2.68             9.0         111    5000.0        21
1         mpfi  3.47   2.68             9.0         111    5000.0        21
2         mpfi  2.68   3.47             9.0         154    5000.0        19
3         mpfi  3.19   3.40            10.0         102    5500.0        24
4         mpfi  3.19   3.40             8.0         115    5500.0        18

   high-L  price  city-L/100km
0      27  13495.0    11.190476
1      27  16500.0    11.190476
2      26  16500.0    12.368421
3      30  13950.0     9.791667
4      22  17450.0    13.055556
```

[5 rows x 27 columns]

[Double-click here for the solution.](#)

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

Example

To demonstrate normalization, let’s say we want to scale the columns “length”, “width” and

“height”

Target: would like to Normalize those variables so their value ranges from 0 to 1.

Approach: replace original value by (original value)/(maximum value)

```
[73]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Questiont #3:

According to the example above, normalize the column “height”.

```
[74]: # Write your code below and press Shift+Enter to execute
```

Double-click here for the solution.

Here we can see, we’ve normalized “length”, “width” and “height” in the range of [0,1].

Binning

Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical ‘bins’, for grouped analysis.

Example:

In our dataset, “horsepower” is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three ‘bins’ to simplify analysis?

We will use the Pandas method ‘cut’ to segment the ‘horsepower’ column into 3 bins

Example of Binning Data In Pandas

Convert data to correct format

```
[75]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```


Lets plot the histogram of horsepower, to see what the distribution of horsepower looks like.

```
[76]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```



```
[76]: Text(0.5, 1.0, 'horsepower bins')
```



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
[77]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[77]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
[78]: group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
[79]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
    ↪include_lowest=True )
df[['horsepower', 'horsepower-binned']].head(20)
```

```
[79]:      horsepower horsepower-binned
0          111             Low
1          111             Low
2          154           Medium
3          102             Low
4          115             Low
5          110             Low
6          110             Low
7          110             Low
8          140           Medium
9          101             Low
10         101             Low
11         121           Medium
12         121           Medium
13         121           Medium
14         182           Medium
15         182           Medium
16         182           Medium
17          48             Low
18          70             Low
19          70             Low
```

Lets see the number of vehicles in each bin.

```
[80]: df["horsepower-binned"].value_counts()
```


```
[80]: Low          153
      Medium       43
      High         5
      Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
[81]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[81]: Text(0.5, 1.0, 'horsepower bins')
```



Check the dataframe above carefully, you will find the last column provides the bins for “horsepower” with 3 categories (“Low”, “Medium” and “High”).

We successfully narrow the intervals from 57 to 3!

Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.


```
[82]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot

a = (0,1,2)

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[82]: Text(0.5, 1.0, 'horsepower bins')
```



The plot above shows the binning result for attribute “horsepower”.

Indicator variable (or dummy variable)

What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called ‘dummies’ because the numbers themselves don’t have inherent meaning.

Why we use indicator variables?

So we can use categorical variables for regression analysis in the later modules.

Example

We see the column “fuel-type” has two unique values, “gas” or “diesel”. Regression doesn’t understand words, only numbers. To use this attribute in regression analysis, we convert “fuel-type” into indicator variables.

We will use the panda’s method ‘get_dummies’ to assign numerical values to different categories of fuel type.

```
[83]: df.columns
```

```
[83]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
        'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
        'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
        'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
```

```
'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'high-L',
'price', 'city-L/100km', 'horsepower-binned'],
dtype='object')
```

get indicator variables and assign it to data frame “dummy_variable_1”

```
[84]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
dummy_variable_1.head()
```

```
[84]:   diesel  gas
0       0    1
1       0    1
2       0    1
3       0    1
4       0    1
```

change column names for clarity

```
[85]: dummy_variable_1.rename(columns={'fuel-type-diesel':'gas', 'fuel-type-gas':
↳ 'diesel'}, inplace=True)
dummy_variable_1.head()
```

```
[85]:   diesel  gas
0       0    1
1       0    1
2       0    1
3       0    1
4       0    1
```

We now have the value 0 to represent “gas” and 1 to represent “diesel” in the column “fuel-type”. We will now insert this column back into our original dataset.

```
[86]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)
```

```
[87]: df.head()
```

```
[87]:   symboling  normalized-losses      make aspiration num-of-doors \
0         3           122  alfa-romero      std           two
1         3           122  alfa-romero      std           two
2         1           122  alfa-romero      std           two
3         2           164      audi      std           four
4         2           164      audi      std           four

   body-style drive-wheels engine-location  wheel-base  length  ... \
```

0	convertible	rwd	front	88.6	0.811148	...
1	convertible	rwd	front	88.6	0.811148	...
2	hatchback	rwd	front	94.5	0.822681	...
3	sedan	fwd	front	99.8	0.848630	...
4	sedan	4wd	front	99.4	0.848630	...

	compression-ratio	horsepower	peak-rpm	city-mpg	high-L	price \
0	9.0	111	5000.0	21	27	13495.0
1	9.0	111	5000.0	21	27	16500.0
2	9.0	154	5000.0	19	26	16500.0
3	10.0	102	5500.0	24	30	13950.0
4	8.0	115	5500.0	18	22	17450.0

	city-L/100km	horsepower-binned	diesel	gas
0	11.190476	Low	0	1
1	11.190476	Low	0	1
2	12.368421	Medium	0	1
3	9.791667	Low	0	1
4	13.055556	Low	0	1

[5 rows x 29 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

Question #4:

As above, create indicator variable to the column of "aspiration": "std" to 0, while "turbo" to 1.

[88]: *# Write your code below and press Shift+Enter to execute*

Double-click here for the solution.

Question #5:

Merge the new dataframe to the original dataframe then drop the column 'aspiration'

[89]: *# Write your code below and press Shift+Enter to execute*

Double-click here for the solution.

save the new csv

[90]: `df.to_csv('clean_df.csv')`

Thank you for completing this notebook

<p><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.